

# SQL Server : Scale out & performance

Christophe Laporte

#SQLSatMadrid



#568 | MADRID 2016

# Apologies – French speaker

---



# Christophe Laporte



~ since 1997  
6.5 <= SQL Server <= 2016



christophe\_laporte@hotmail.fr



<http://conseilit.wordpress.com/>



#568 | MADRID 2016



# BIG Thanks to SQLSatMadrid Sponsors

---



**Hewlett Packard  
Enterprise**



**plain  
concepts**



#SQLSatMadrid



#568 | MADRID 2016



## 4 Sponsor Sessions at 11:40

---

Don't miss them, they might be getting distributing some awesome prizes!

- HPE
- SolidQ
- KABEL
- TSD Consulting

Also **BIG Raffle prizes** at the end of the event provided by:

Plainconcepts, SolidQ, Kabel, TSD Consulting, Pyramid Analytics & sqlpass.es


# Agenda

---

- Scaling for performance

# What is scaling { up | out }

- *Scalability* is the ability of an application to efficiently use more resources in order to do more useful work.
- Scale Up
  - Bigger box
- Scale Out
  - Expending to multiple servers
- Main idea
  - Run queries faster
  - Run more queries at the same time



# Scale up

---

- Historically
  - Processors speed
  - More cores
  - More RAM
  - Disks faster and faster (SSD/PCI/NVMe and NVDIMM)
- Problem with scaling up
  - Fault tolerance very expensive
  - Scaling on-demand
    - Support high-end workload
    - Workload changes



# But ...

---

- Locks & Latches
  - RCSI
  - In memory tables
- Only 1 TempDB
  - Adding more files to decrease contention (GAM / SGAM / PFS)
  - SQL 2014 : TempDB eager write
  - SQL 2016 : new Query hint NO\_PERFORMANCE\_SPOOL
- 1 transaction log file / Limit of Log Manager
  - Delayed durability
  - Schema only durability for Hekaton tables

# Scale out

---

- Scale out - the future ?
  - SQL
  - NoSQL
- Scale out often means fault tolerant
- Small DBs / instances
  - Easier to manage
  - Lower cost / commodity hardware
  - Geographically dispersed
    - Data nearest of the end-user
    - Network latency

# How to scale out ? (the real agenda !)


---

- Middle Tier
- Distributed partitioned view
- Replication
- Log shipping
- Scalable Shared Databases
- Partitioned tables
- Availability groups
- Service broker

# Middle Tier/ data layer

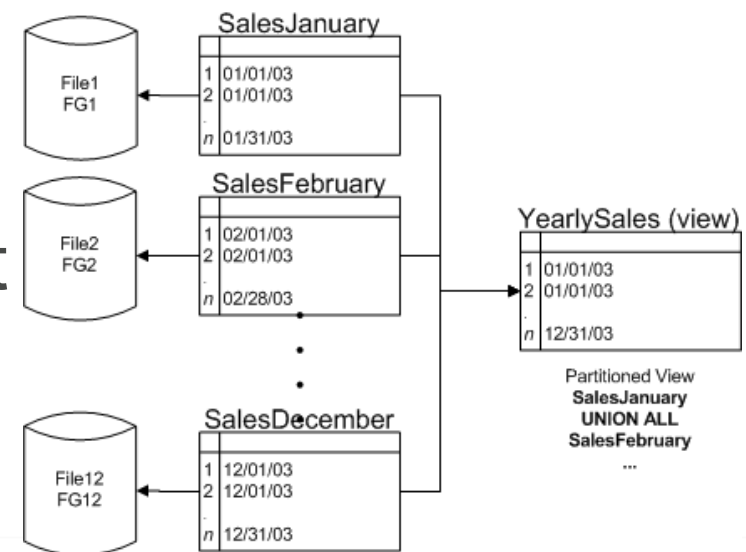
---

- Data-dependent routing from middleware
- Table / database / instance chosen by code
- Boost performance with caching
- Can be combined



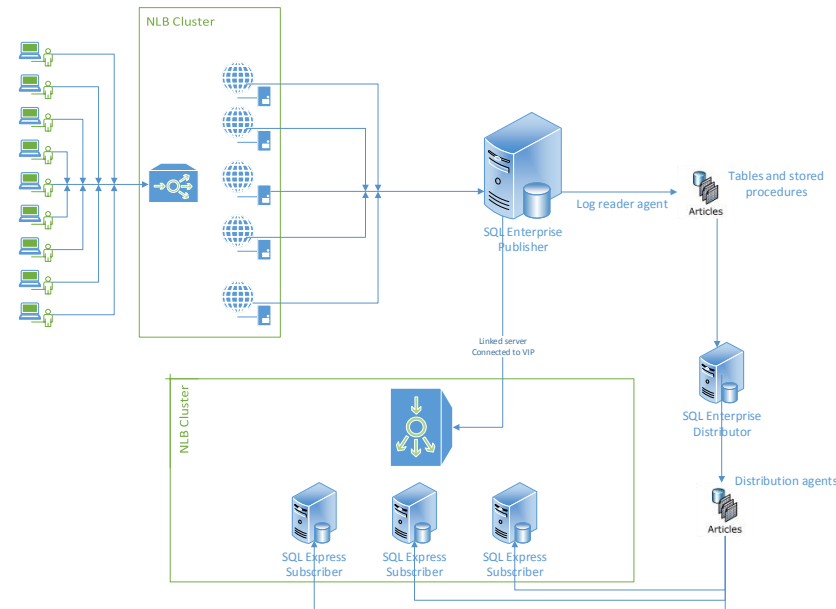
# (Distributed) Partitioned Views

- Can be spread across multiple servers
- Local joins to not move data across servers
- Smaller tables : easy to manage (Index, stats, checkdb)
- Horizontal partitioning
- Take care to restore point
- Old fashion (SQL 2000)
- Demo 1



# Replication


- Read-only
  - Easy to setup and manage
  - 1 publisher / N subscribers
  - N publishers / 1 subscriber
  - transactional replication
- Read-Write
  - Merge replication
  - 1..N publishers / 1..N subscribers
  - Slightly more complex to setup and manage
- Require additional feature to load balance






# Peer-To-Peer Replication

- Enterprise edition
- May be geographically dispersed
- Gaps / identity columns
- Update frequency moderate




# Log Shipping

- Read-only
- all editions
- Large # of secondary
- Latency to get up-to-date data
- May be geographically dispersed
- Users disconnected during restores
- Require additional feature to load balance




# Scalable Shared Databases

- Enterprise edition
- Read-only
  - Updating the DB may be difficult / decrease availability
- Require additional feature to load balance



# Partitioned Tables

- Lock reduction (lock escalation at partition level)
- Deal with different kind of storage / compression level
- Truncate partition (new in SQL 2016)
- Resolves page latch contention issues
- Scaling out or scaling up ?
- Edition enterprise
- Demo 2



# Availability Groups – RO load Balancing

- Read Only Load Balancing : New in SQL Server 2016
- Easy setup & does not require high skills on system administration
- Domain agnostic
- Up to 8 RO replicas
- 3 synchronous
- Enterprise edition
- Demo 3

✓ DemoAG: hosted by SQL2K16W01 (Replica role: Primary) Last updated: 08/06/2016 22:27:46  
Auto refresh: on

Availability group state: ✓ Healthy  
Primary instance: SQL2K16W01  
Failover mode: Automatic  
Cluster state: CLUSTSQLWRK (Normal Quorum)  
[Start Failover Wizard](#)  
[View AlwaysOn Health Events](#)  
[View Cluster Quorum Information](#)

Availability replica: [Add/Remove Columns](#)


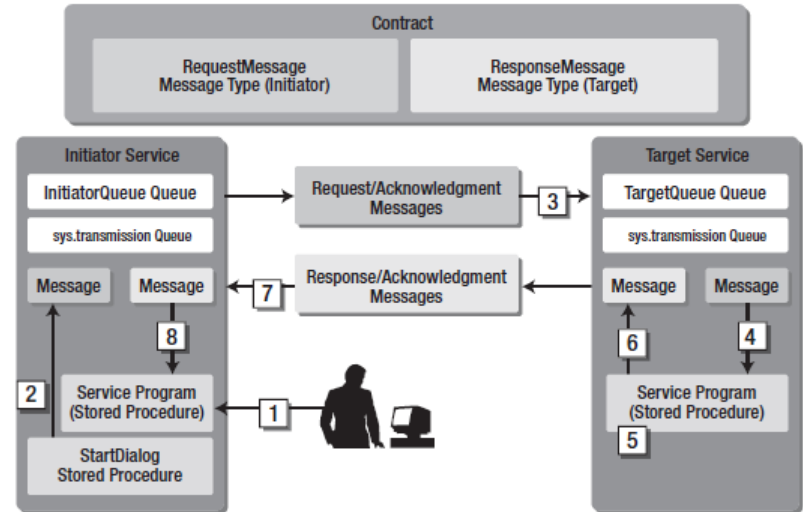
Name	Role	Failover Mode	Synchronization State	Issues
✓ <a href="#">SQL2K16W01</a>	Primary	Automatic	Synchronized	
✓ <a href="#">SQL2K16W02</a>	Secon...	Automatic	Synchronized	
✓ <a href="#">SQL2K16W03</a>	Secon...	Automatic	Synchronized	

Group by ▾ [Add/Remove Columns](#)

Name	Replica	Synchronization State	Failover Read...	Issues
<a href="#">SQL2K16W01</a>				^
✓ DemoDB	SQL2K16W01	Synchronized	No Data Loss	
<a href="#">SQL2K16W02</a>				^
✓ DemoDB	SQL2K16W02	Synchronized	No Data Loss	
<a href="#">SQL2K16W03</a>				^
✓ DemoDB	SQL2K16W03	Synchronized	No Data Loss	

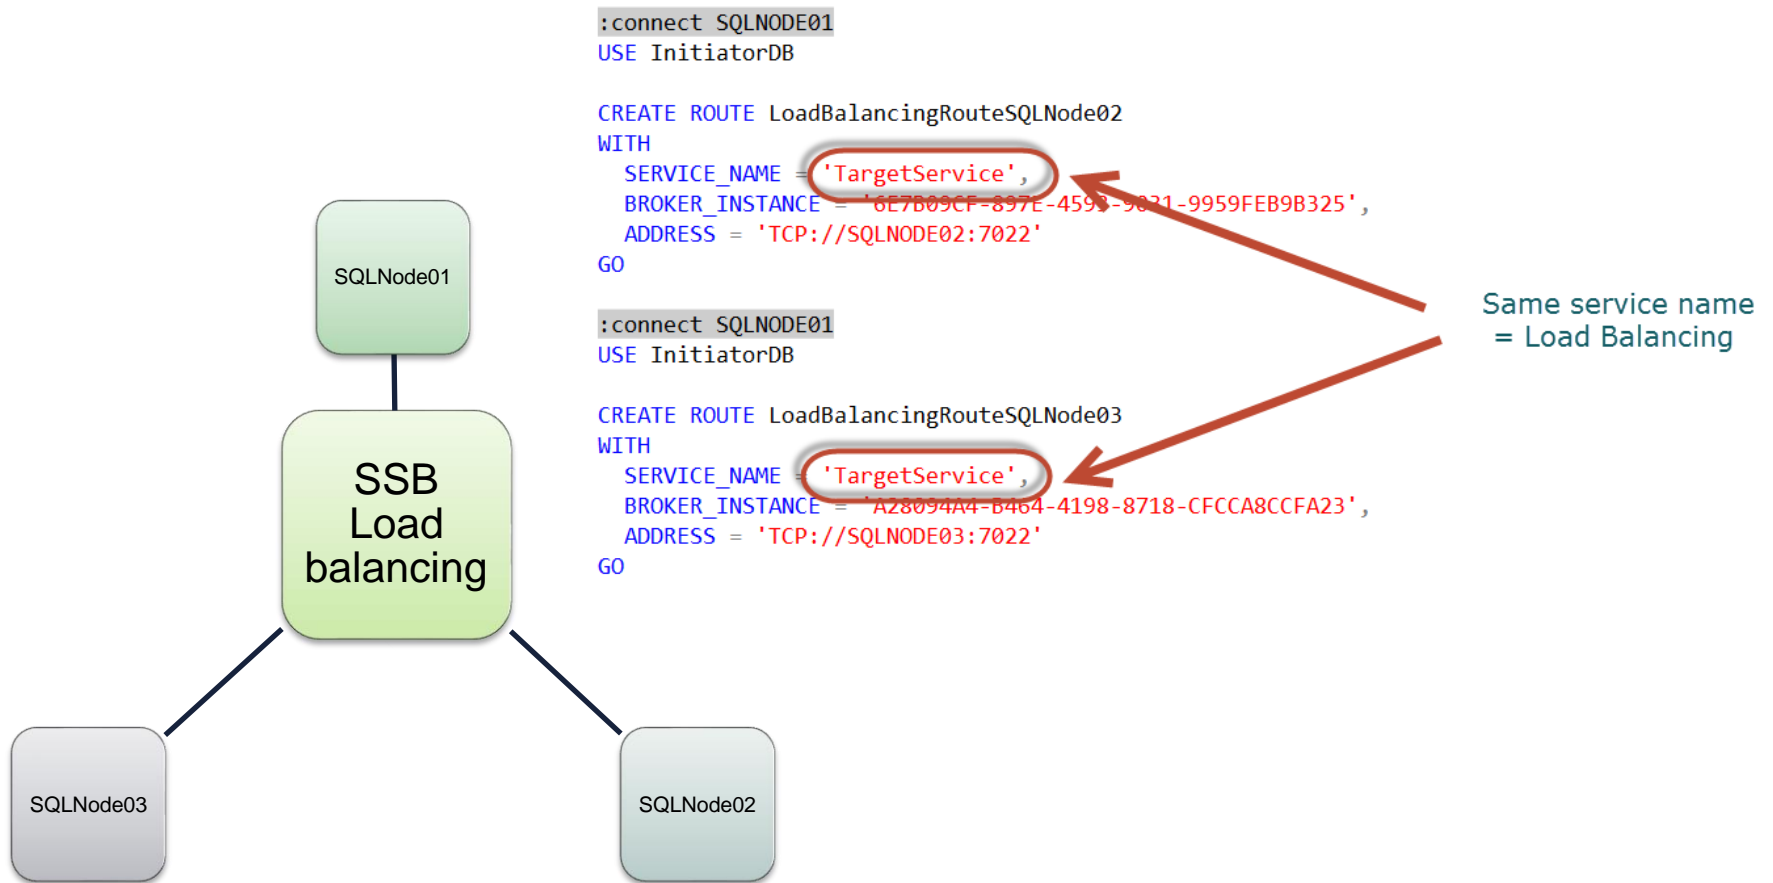
# Service Broker

- Asynchronous operations
- Sends message to a target
  - Same database
  - Same instance different database
  - Different instance
- Routing between instances
  - Load balancing
  - Multicasting



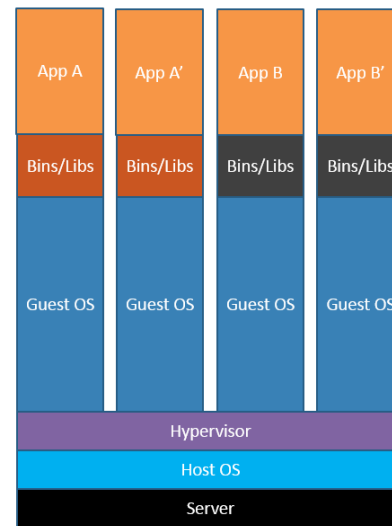



# Demo 4 – Service broker & NLB




# Windows Server containers / Docker

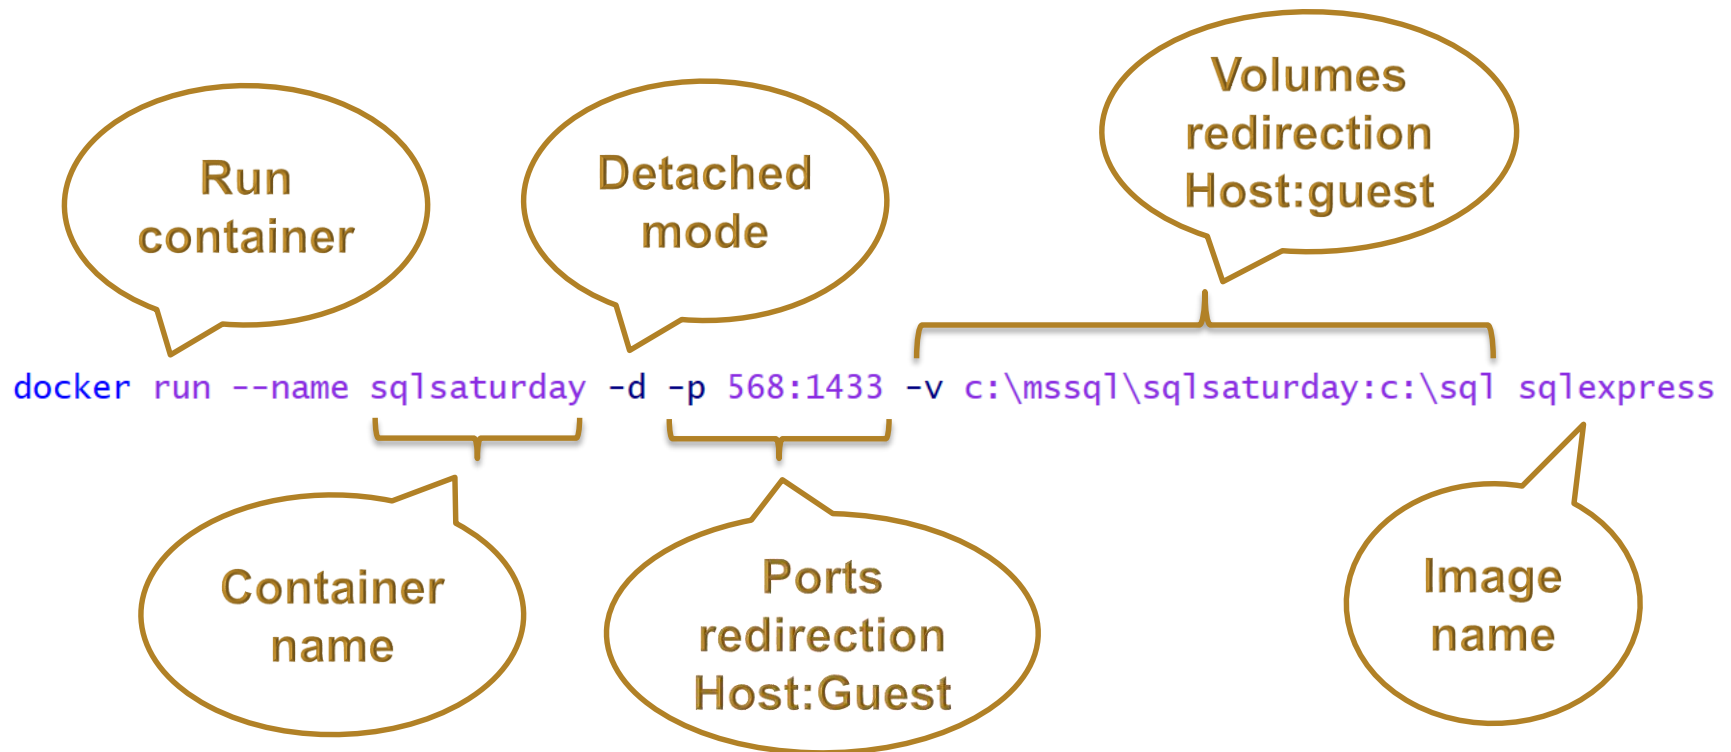
- Use cases : Dev / Test / Prod
- Lightweight SQL Server instance inside container
- Quick response to load increase
- Scaling out if combined with other solutions
  - Third party load balancer
  - (/\ Replication /\ )
  - Service Broker



Virtual machines  
versus containers



# IT Demo – SQL Server Express in a container



```
PS C:\Users\Administrator> docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
aea5623fd5cd	sql express	"cmd /S /C powershell"	3 days ago	Up 3 days	0.0.0.0:61439->1433/tcp, 0.0.0.0:61779->7022/tcp	sql09
dd76aa08926c	sql express	"cmd /S /C powershell"	3 days ago	Up 3 days	0.0.0.0:61438->1433/tcp, 0.0.0.0:61778->7022/tcp	sql08
bab00c25e867	sql express	"cmd /S /C powershell"	3 days ago	Up 3 days	0.0.0.0:61437->1433/tcp, 0.0.0.0:61777->7022/tcp	sql07
3a3299f4e24d	sql express	"cmd /S /C powershell"	3 days ago	Up 3 days	0.0.0.0:61435->1433/tcp, 0.0.0.0:61775->7022/tcp	sql05
860b0f630541	sql express	"cmd /S /C powershell"	3 days ago	Up 3 days	0.0.0.0:61434->1433/tcp, 0.0.0.0:61774->7022/tcp	sql04
e0d46a11f1ff	sql express	"cmd /S /C powershell"	3 days ago	Up 3 days	0.0.0.0:61433->1433/tcp, 0.0.0.0:61773->7022/tcp	sql03
031e615766a7	sql express	"cmd /S /C powershell"	3 days ago	Up 3 days	0.0.0.0:61432->1433/tcp, 0.0.0.0:61772->7022/tcp	sql02
bc0da018328e	sql express	"cmd /S /C powershell"	3 days ago	Up 3 days	0.0.0.0:61431->1433/tcp, 0.0.0.0:61771->7022/tcp	sql01

#SQLSatMadrid



# Demo 5 – Service Broker & Multicasting with Docker

```
BEGIN TRANSACTION;
```

```
BEGIN DIALOG @DialogHandleSQL01  
FROM SERVICE InitiatorService  
TO SERVICE N'TargetServiceSQL01'  
ON CONTRACT SampleContract  
WITH ENCRYPTION = OFF;
```

```
BEGIN DIALOG @DialogHandleSQL02  
FROM SERVICE InitiatorService  
TO SERVICE N'TargetServiceSQL02'  
ON CONTRACT SampleContract  
WITH ENCRYPTION = OFF;
```

```
BEGIN DIALOG @DialogHandleSQL03  
FROM SERVICE InitiatorService  
TO SERVICE N'TargetServiceSQL03'  
ON CONTRACT SampleContract  
WITH ENCRYPTION = OFF;
```

```
BEGIN DIALOG @DialogHandleSQL04  
FROM SERVICE InitiatorService  
TO SERVICE N'TargetServiceSQL04'  
ON CONTRACT SampleContract  
WITH ENCRYPTION = OFF;
```

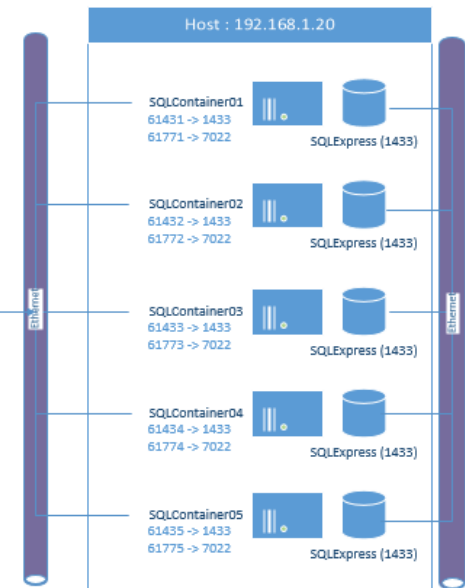

```
BEGIN DIALOG @DialogHandleSQL05  
FROM SERVICE InitiatorService  
TO SERVICE N'TargetServiceSQL05'  
ON CONTRACT SampleContract  
WITH ENCRYPTION = OFF;
```

```
SEND ON CONVERSATION (  
    @DialogHandleSQL01,@DialogHandleSQL02,@DialogHandleSQL03,@DialogHandleSQL04,@DialogHandleSQL05  
)  
MESSAGE TYPE RequestMessage (@RequestMsg);
```

```
SELECT @RequestMsg AS SentRequestMsg;
```

```
COMMIT TRANSACTION;
```

#SQLSatMadrid



Multiple Begin Dialog  
Single Send on Conversation

# How to load balance workload

---

- Windows NLB
- Windows 2016 Software Load Balancing (SLB) for SDN
- 3rd party hardware
  - Kemp
  - F5
  - Cisco
  - Citrix
  - Radware
- 3rd party software
  - Kemp
  - Pfsense
  - Scalearc
  - Nginx
  - ...

# Demo 6 – load balancing with containers

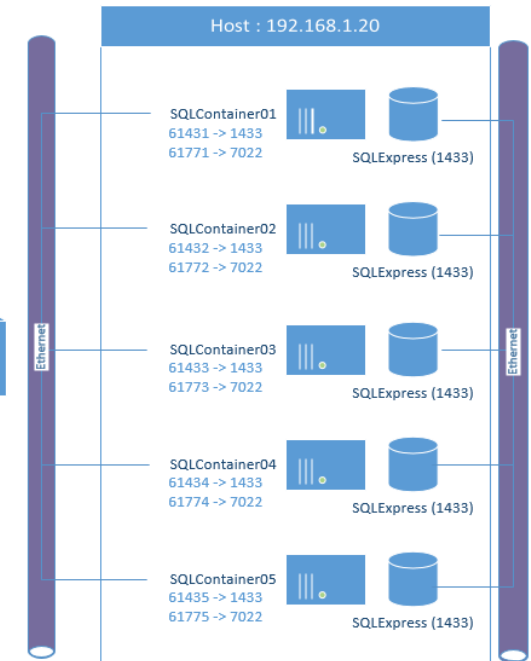
Name	Virtual IP Address	Protocol	Status	Total Conns	Last 60 Sec	5 Mins	30 Mins	1 Hour	Active Conns	Current Rate Conns/s	Real Servers RS-IP	% Conns/s
1 NLBMSSQL	192.168.1.128:1433	tcp	Up	557	100	150	150	150	50	10	192.168.1.20:61431 192.168.1.20:61432 192.168.1.20:61433 192.168.1.20:61434 192.168.1.20:61435	20 20 20 20 20
1	System Total Conns			557	100	150	150	150	50	10 /sec		

@ VIP ,1433



Load balancer

		Queue			Session rate			Sessions					Bytes			
		Cur	Max	Limit	Cur	Max	Limit	Cur	Max	Limit	Total	LbTot	Last	In	Out	
<input type="checkbox"/>	Conteneur01	0	0	-	0	10		0	10	-	10	10	4h21m	9 180	15 780	
<input type="checkbox"/>	Conteneur02	0	0	-	0	10		0	10	-	10	10	4h21m	9 180	15 780	
<input type="checkbox"/>	Conteneur03	0	0	-	0	10		0	10	-	10	10	4h21m	9 180	15 780	
<input type="checkbox"/>	Conteneur04	0	0	-	0	10		0	10	-	10	10	4h21m	9 180	15 780	
<input type="checkbox"/>	Conteneur05	0	0	-	0	10		0	10	-	10	10	4h21m	9 180	15 780	
	Backend	0	0		0	50		0	50	200	50	50	4h21m	45 800	78 900	





# SQL Server on Azure ...

---

- Not covered in this session
- By design
  - SQL DataWarehouse
    - Elastic data warehouse as a service
  - Azure SQL Databases
    - Sharding / Federation
    - Elastic pools
- Scale up/down
  - Change service-tiers for a given database
- Scale out
  - Add Remove databases

#SQLSatMadrid



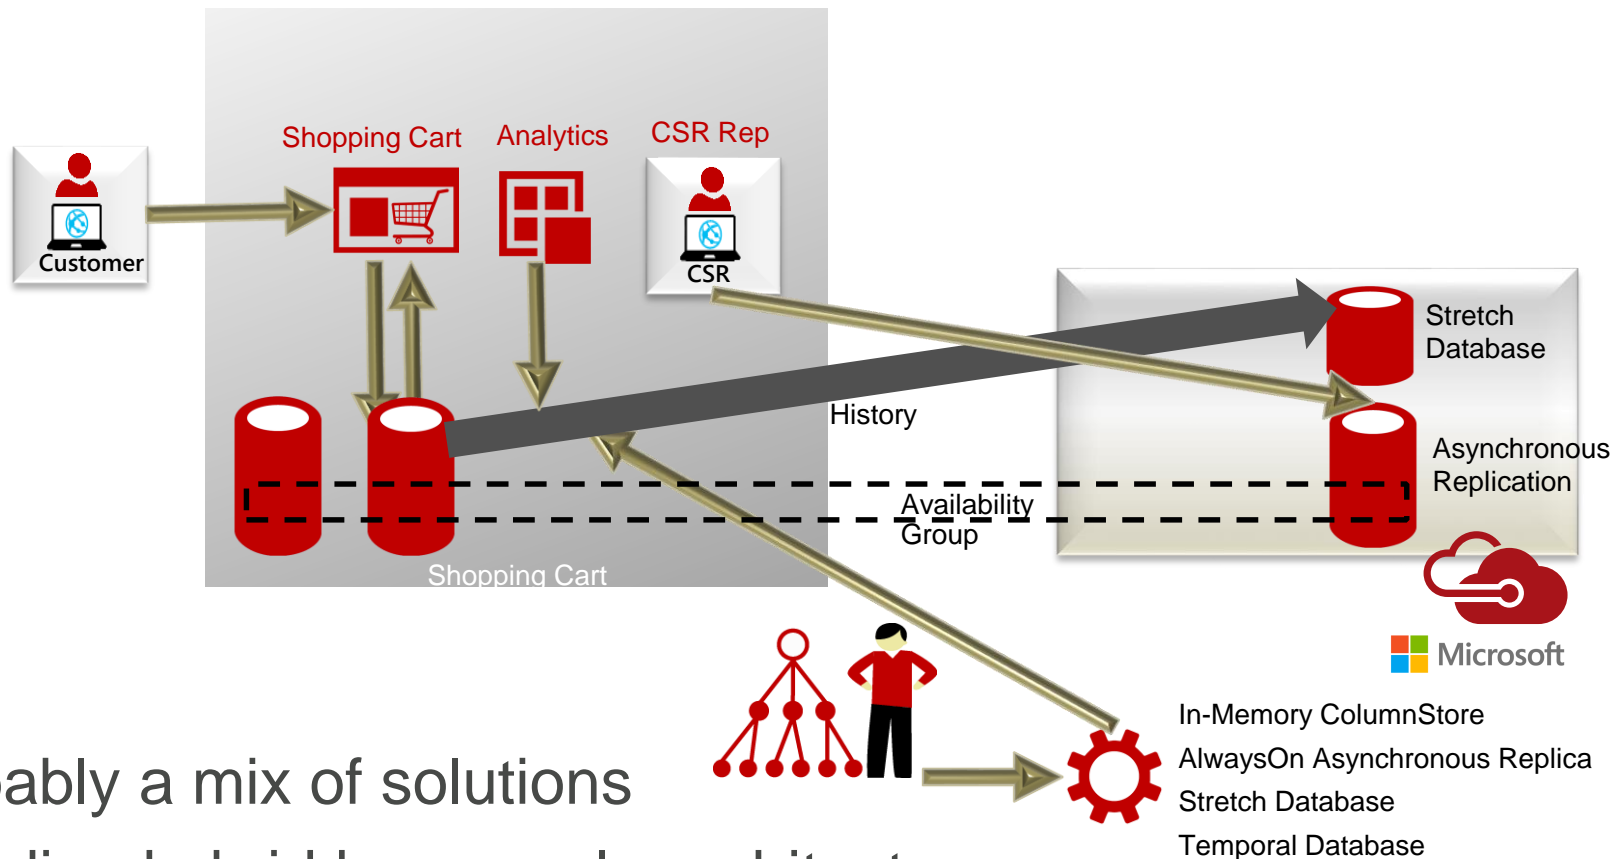
#568 | MADRID 2016

# Scaling : remember some basics

---

- A clustered index on an integer column does not scale !
  - Unless partitioning
- A GUID / uniqueidentifier does
  - But introduce fragmentation

# The future ?



- Probably a mix of solutions
- Including hybrid hyperscale architecture

#SQLSatMadrid

# Q/A

---

- Hope you enjoyed
- Thanks for attending
- Q/A